

---

# COMP3308

# Assignment 1

---

Due: 5:00pm Friday, 8 May 2015

*This assignment is worth 20% of your final mark*

## Task description

In this assignment you will implement the K-Nearest Neighbour and Naïve Bayes algorithms and evaluate them on a real dataset using the stratified cross validation method. You will also evaluate the performance of other classifiers on the same dataset using Weka. Finally, you will investigate the effect of feature selection, in particular the Correlation-based Feature Selection method (CFS) from Weka.

## Late submissions policy

Penalty for late submission is 1 mark per day, including the weekends. If the assignment is submitted more than 7 days after the deadline, it will not be accepted. The day is considered to end at 5:00pm, i.e. if the assignment is submitted at 5:05pm on 8th of May it will be considered one day late.

## Programming languages

Your implementation can be written in Python, Java, C, C++ or Octave. The assignment will be tested on the University machines, so your code must be compatible with the language version installed on those machines. You are not allowed to use any of the built-in classification libraries for the purposes of this assignment.

**Note:** If you are familiar with Matlab, have a look at this wiki entry on the differences between Matlab and Octave: [http://wiki.octave.org/FAQ#Porting\\_programs\\_from\\_Matlab\\_to\\_Octave](http://wiki.octave.org/FAQ#Porting_programs_from_Matlab_to_Octave). Licensing issues require us to run Octave on the submission servers rather than Matlab, so make sure your submission is compatible with Octave.

## Submission and pair work

This assignment will be submitted using the submission tool PASTA (<http://comp3308.it.usyd.edu.au>). In order to connect to the website, you'll need to be connected to the university VPN. You can read [this page](#) to find out how to connect to the VPN.

Your assignment can be completed individually or in pairs. See the submission details section for more information about how to submit.

PASTA will allow you to make as many submissions as you wish, and each submission will provide you with feedback on each of the components of the assignment. Your last submission before the assignment deadline will be marked, and the mark displayed on PASTA will be the final mark for your code (12 marks).

## 1 Data

The dataset for this assignment is the Pima Indian Diabetes dataset. It contains 768 instances described by 8 numeric attributes. There are two classes - **yes** and **no**. Each entry in the dataset corresponds to a patient's record; the attributes are personal characteristics and test measurements; the class shows if the person shows signs of diabetes or not. The patients are from Pima Indian heritage, hence the name of the dataset.

A copy of the dataset can be downloaded from Blackboard. There are 2 files associated with the dataset. The first file, **\*.names**, describes the data, including the number and the type of the attributes and classes, as well as their meaning. The second file, **\*.data**, contains the data itself. Your task is to predict the class, where the class can be **yes** or **no**.

**Note:** The original dataset can be sourced from UCI Machine Learning Repository. However, you need to use the dataset available on Blackboard as it has been modified for consistency.

## 2 Data preprocessing

Read the **pima-indians-diabetes.names** file and learn more about the meaning of the attributes and the classes. Use Weka's in-built normalisation filter to normalise the values of each attribute to make sure they are in the range [0,1]. The normalisation should be done along each column (attribute), not each row (entry). The class attribute is not normalised – it should remain unchanged. Save the preprocessed file as **pima.csv**.

**Warning:** In order to ensure that Weka can process the data, you will need to add headers to the data file and save it as a .csv file. The headers should be removed after preprocessing.

## 3 Classification algorithms

### K-Nearest Neighbour

The K-Nearest Neighbour algorithm should be implemented for any K value and should use Euclidean distance as the distance measure. If there is ever a tie between the two classes, choose class **yes**.

### Naïve Bayes

The Naïve Bayes should be implemented for numeric attributes, using a probability density function. Assume a normal distribution, i.e. use the probability density function for a normal distribution. As before, if there is ever a tie between the two classes, choose class **yes**.

**Note:** Carefully read section 6 to find out how your program will be expected to receive input and give output.

## 4 10-fold stratified cross-validation

In order to evaluate the performance of the classifiers, you will have to implement 10-fold stratified cross-validation. Your program should be able to show the algorithm's average accuracy over the 10 folds. This information will be required to complete the report.

Your implementation of 10-fold stratified cross-validation will be tested based on your `pima-folds.csv` file. The information about the folds should be stored in `pima-folds.csv` in the following format for each fold:

- Name of the fold, `fold1` to `fold10`.
- Contents of the fold, with each entry on a new line.
- A single blank line to separate the folds from each other.

An example of the `pima-folds.csv` file would look as follows:

```
fold1
0.588,0.628,0.574,0.263,0.136,0.463,0.054,0.333,yes
0.243,0.274,0.224,0.894,0.113,0.168,0.735,0.321,no

fold2
0.588,0.628,0.574,0.263,0.136,0.463,0.054,0.333,yes
0.243,0.274,0.224,0.894,0.113,0.168,0.735,0.321,no

...
fold10
0.588,0.628,0.574,0.263,0.136,0.463,0.054,0.333,yes
0.243,0.274,0.224,0.894,0.113,0.168,0.735,0.321,no
```

**Note:** The number of instances per fold should not vary by more than one. If the total number of instances is not divisible by ten, the remaining items should be distributed amongst the folds rather than being placed in one fold.

## 5 Feature selection

Correlation-based feature selection (CFS) is a method for selecting a subset of the original attributes. It searches for the best subset of features, where best is defined by a heuristic which considers how good the individual features are at predicting the class and how much they correlate with the other features. Good subsets of features contain features that are highly correlated with the class and uncorrelated with each other.

Load the `pima.csv` file in Weka, and apply CFS to reduce the number of features. It is available from the “Select attributes” tab in Weka. Use “Best-First Search” as the search method. Save the CSV file with the reduced number of attributes (this can be done in Weka) and name it `pima-CFS.csv`.

**Warning:** As before, in order to ensure Weka can understand the data, you'll need to add headers. Once you are done processing, remove the headers

## 6 Input and output

### Input

Your program will need to be named **MyProgram**, however may be written in any of the languages mentioned in the “Programming languages” section.

Your program should take 3 command line arguments. The first argument is the path to the training data file, the second is the path to the testing data file, and the third is the name of the algorithm to be executed (**NB** for Naïve Bayes and **kNN** for the Nearest Neighbour, where **k** is replaced with a number; e.g. **5NN**).

For example, if you were to make a submission in Java, your main class would be **MyProgram.java**, and the following are possible inputs to the program:

```
$ java MyProgram pima.csv examples.csv NB
$ java MyProgram pima-CFS.csv examples.csv 4NN
```

The input testing data file will consist of several new examples to test your data on. The file will not have headers, will have one example per line, and each line will consist of a normalised value for each of the 8 non-class attributes separated by commas. An example input file would look as follows:

```
0.588,0.628,0.574,0.263,0.136,0.463,0.054,0.333
0.243,0.274,0.224,0.894,0.113,0.168,0.735,0.321
0.738,0.295,0.924,0.113,0.693,0.666,0.486,0.525
```

### Output

Your program will output to standard output (a.k.a. “the console”). The output should be one class value (**yes** or **no**) per line – each line representing your program’s classification of the corresponding line in the input file. An example output should look as follows:

```
yes
no
yes
```

**Note:** If you have any questions or need any clarifications about program input or output, ask a question on Piazza or ask your tutor. Since your program will be automatically tested by PASTA, it is important that you follow the instructions exactly.

## 7 Weka evaluation

In Weka select 10-fold cross validation (it is actually 10-fold stratified cross validation) and run the following algorithms: ZeroR, 1R, k-Nearest Neighbor (k-NN; IBk in Weka), Naïve Bayes (NB), Decision Tree (DT; J48 in Weka) and Multi-Layer Perceptron (MLP).

Compare their performance with your k-Nearest Neighbor and Naïve Bayes classifiers. Do this for the case without feature selection (using **pima.csv**) and with CFS feature selection (using **pima-CFS.csv**).

## 8 Report

You will have to describe your analysis and findings in a report similar to a research paper. Your report should include 5 sections. There is no minimum or maximum length for the report – you will be marked on the quality of the content that you provide.

### 8.1 Aim

This section should briefly state the aim of your study, e.g. predicting X based on Y etc., and include a paragraph about why this problem is important.

### 8.2 Data

This section should describe the dataset, mentioning the number of attributes and classes. It should also briefly describe the CFS method and list the attributes selected by the CFS.

### 8.3 Results and discussion

The accuracy results should be presented (in percentage, using 10-fold cross validation) in the following table where My1NN, My5NN and MyNB are your implementations of the 1NN, 5NN and NB algorithms, evaluated using your stratified 10-fold cross validation.

	ZeroR	1R	1NN	5NN	NB	DT	MLP
No feature selection							
CFS							

	My1NN	My5NN	MyNB
No feature selection			
CFS			

In the discussion, compare the performance of the classifiers, with and without feature selection. Compare your implementations of k-NN and NB with Weka's. Discuss the effect of the feature selection – did CFS select a subset of the original features, and if so, did the selected subset make intuitive sense to you? Was feature selection beneficial, i.e. did it improve accuracy, or have any other advantages? Why do you think this is the case? Include anything else that you consider important.

### 8.4 Conclusion

Summarise your main findings and, if possible, suggest future work.

### 8.5 Reflection

Write one or two paragraphs describing the most important thing that you have learned throughout this assignment.

## Submission details

This assignment is to be submitted electronically via the PASTA submission system.

### Individual submissions

If you are working individually, your submission should be zipped together in a single **.zip** file and include the following:

- The report in PDF format.
- The source code with a main program called **MyProgram**. Valid extensions are **.java**, **.py**, **.c**, **.cpp**, **.cc**, and **.m**.
- Three data files: **pima.csv**, **pima-CFS.csv** and **pima-folds.csv**.

A valid submission might look like this:

```
submission.zip
|- pima.csv
|- pima-folds.csv
|- pima-CFS.csv
|- report/
|  +- report.pdf
|- MyProgram.java
+- extrapackage/
   |- MyClass.java
   +- OtherClass.java
```

Upload your submission on PASTA under **Assignment 1**.

### Pair submissions

If you are working in a pair, only one person can submit the assignment. If both group members submit accidentally, one should submit an empty zip file and notify their tutor.

The group member who chooses to submit should submit the pair's work according to the individual assignment instructions above.

When submitting group work, you are also required to make a group work submission. This will be zipped into a single **.zip** file and include the following:

- Filled in group work assignment cover sheet in PDF format.
- A file name **group.txt** that contains both of group members' unikeys. Each unikey should be on a new line.

The same group member who previously submitted should then upload this file on PASTA under **Assignment 1 Pair Coversheet**.

**Note:** The group work cover sheet can be downloaded at [http://sydney.edu.au/engineering/it/current\\_students/postgrad\\_coursework/policies/academic\\_honesty.shtml](http://sydney.edu.au/engineering/it/current_students/postgrad_coursework/policies/academic_honesty.shtml).

## Marking criteria

	Your mark	Marker's comments
[12 marks] Code – based on the tests in PASTA; automatic marking	see PASTA	see PASTA
<p>[8 marks] Report:</p> <p>[0.5 marks] Introduction</p> <ul style="list-style-type: none"> <li>• What is the aim of the study?</li> <li>• Why is this study (the problem) important?</li> </ul> <p>[0.5 marks] Data – well explained</p> <ul style="list-style-type: none"> <li>• Dataset – brief description of the dataset</li> <li>• Attribute selection – brief summary of CFS and stating the selected attributes</li> </ul> <p>[4 marks] Results and discussion</p> <ul style="list-style-type: none"> <li>• All results presented</li> <li>• Correct and deep discussion of the results</li> <li>• Effect of the feature selection – beneficial or not (accuracy, other advantages)</li> <li>• Comparison between the classifiers (accuracy, other advantages)</li> </ul> <p>[1.5 marks] Conclusions and future work</p> <ul style="list-style-type: none"> <li>• Meaningful conclusions based on the results</li> <li>• Meaningful future work suggested</li> </ul> <p>[0.5 marks] Reflection (meaningful and relevant personal reflection)</p> <p>[1 marks] English and presentation</p> <ul style="list-style-type: none"> <li>• Academic style, grammatical sentences, no spelling mistakes</li> <li>• Good structure and layout; consistent formatting</li> </ul>		
<p>Penalty for late submission</p> <ul style="list-style-type: none"> <li>• – 1 mark for each day late</li> </ul>		
Total (out of 20)		